

# A $\text{\LaTeX}$ 'ed solution to B-tree insertion and deletion.

Mitchell Scott (mtscot4)

September 27, 2024

# B-Tree Insertion and Deletion Assignments

Assignment 3-1: Insert


Assignment 3-1: Deletion

# Table of Contents


Assignment 3-1: Insert

Assignment 3-1: Deletion


For  $m = 2$ , and the below B-tree, insert the keys 25, 17, 42 and 29, in this order.




25 is inserted into currentNode, currentNode is checked for overflow. Since  $4 \leq 2m = 4$ , there is no overflow.




17 is inserted into currentNode, currentNode is checked for overflow.  $5 > 2m = 4$ , so overflow, as it violates the B-tree property.



16 is found to be the median of the keys of currentNode, so we move 13 from currentNode to Parent.




Now, parent has the wrong number of pointers, as four keys should have five pointers. We split 'currentNode' into 'leftNode' and 'rightNode'.







We add the left and right pointer to `leftNode` and `rightNode`, respectively. Lastly, since the root has grown, we check for overflow. Since  $4 \leq 2m = 4$  keys, there is no overflow. We have finally inserted 17.




Now we seek to insert 42. Since the new `currentNode` has  $4 \leq 2m = 4$  keys, we don't have overflow, so we are done.



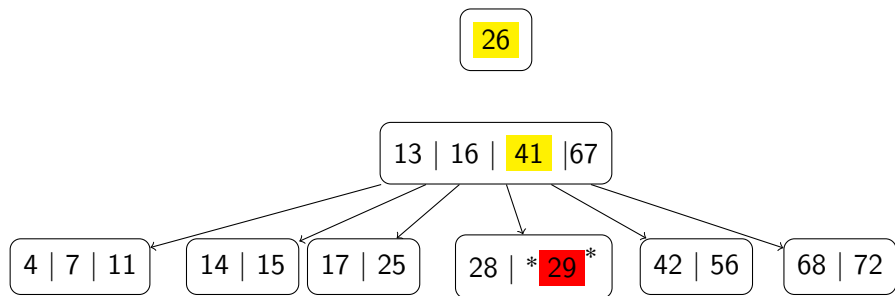
Lastly, we seek to add 29. We insert into `currentNode`, but observe that `currentNode` is overflowing as it has  $5 > 2m = 4$  keys.




41 is the median of `currentNode` so we move it to parent. The rest of the node is split as we need another pointer coming from parent.



Lastly, we have overflow at the root, so we follow the same procedure, where we find the median, and move to parent, or becoming the new root in this case.



We split the remaining node, and they are now children of newRoot. A keen observer, such as RANDOM CAT, might suspect underflow at the new root, but the root doesn't have a minimum number of keys, so we are done!




# Table of Contents

Assignment 3-1: Insert


Assignment 3-1: Deletion

For  $m = 2$ , and the below B-tree, delete the keys 18,14,21, and 1, in this order.







Since 18 is a leaf node, it is deleted at `currentNode`. Then we check `currentNode` for underflow. Since  $3 > m = 2$ , there is no underflow, so no additional work to delete 18.




14 is deleted at currentNode, so we check currentNode for underflow.




currentNode is a leaf node that is underflowing as  $1 < m = 2$  key. We observe that rightSibling has  $3 \geq m = 2$  nodes, so we borrow from it.




We perform a rotation to borrow from `rightSibling`. 17 goes from `rightSibling` to parent, which allows 16 to go from parent to `currentNode`. Done.




To delete 21, which is a branch node, we must switch it with the next largest key, 22, which is guaranteed to be a leaf node.




We now delete key 21, which causes an underflow in `currentNode`. Since `currentNode` has no right sibling, we turn to `leftSibling` to attempt to balance.



Observing that `leftSibling` has exactly  $m = 2$  keys, we cannot rotate, so we construct a `newNode` with exactly  $2m = 4$  keys by collapsing `currentNode` and `leftSibling`.




We then delete `currentNode`, delete `leftSibling` and delete the separator from parent. Then we link the gap in parent to `newNode`.







Lastly, we check for underflow at parent, but since  $2 \geq m = 2$ , there is no underflow, and we are done.




Next we delete key 1, which we can do as it is a leaf node. This leaves underflow at `currentNode`. We look to `rightSibling` and see that it doesn't have more than  $m = 2$  keys.




Since a rotation is not possible, we collapse `currentNode` and `rightSibling`.



When we are done linking the gap from `newNode` to parent, we then check overflow for parent, which it unfortunately is underflowing.



Since we are underflowing at parent, and `rightSibling` cannot help, we must collapse parent and `rightSibling`, to finally arrive at the finished B-tree.



Lastly, we perform a sanity check to see that every node has between 2 and  $2m = 4$  keys, and the parent has 4 nodes and 5 children, so we are done.

